



南開大學
Nankai University

计算机学院
并行程序设计第 4 次作业

高斯消去法的 Pthreads 并行化

姓名：丁屹
学号：2013280
专业：计算机科学与技术

2022 年 5 月 4 日

目录

1 问题描述	2
2 Pthreads 算法设计	3
2.1 测试用例的确定	3
2.2 实验环境和相关配置	3
2.3 算法设计	3
2.3.1 默认平凡算法	3
2.3.2 使用 Pthreads 动态创建线程并行化加速	4
2.3.3 使用 Pthreads 线程池和信号量同步并行化加速	5
2.3.4 使用 Pthreads 线程池和 barrier 栅栏同步并行化加速	9
3 实验结果分析	11

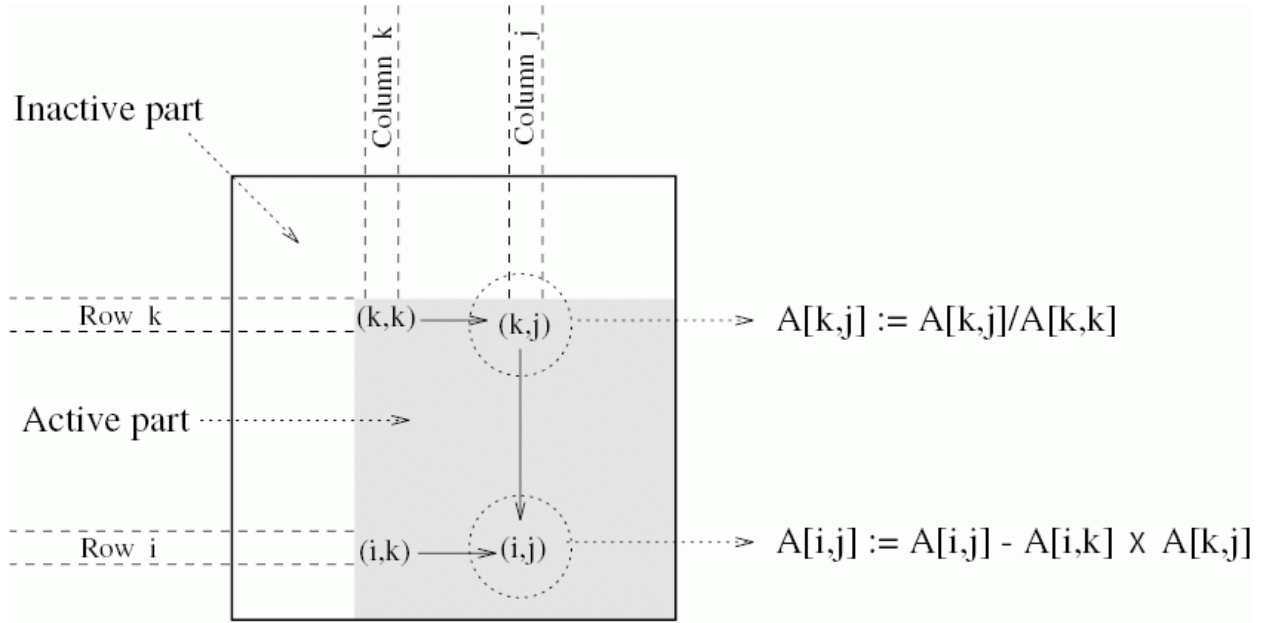


图 1.1: 高斯消去法示意图

1 问题描述

高斯消去的计算模式如图 1.1 所示，在第 k 步时，对第 k 行从 (k,k) 开始进行除法操作，并且将后续的 $k+1$ 至 N 行进行减去第 k 行的操作，串行算法如下面伪代码所示。

Algorithm 1 普通高斯消元算法伪代码

```

1: function LU
2:   for  $k := 0$  to  $n$  do
3:     for  $j := k + 1$  to  $n$  do
4:        $A[k, j] := A[k, j]/A[k, k]$ 
5:     end for
6:      $A[k, k] := 1.0$ 
7:     for  $i := k + 1$  to  $n$  do
8:       for  $j := k + 1$  to  $n$  do
9:          $A[i, j] := A[i, j] - A[i, k] * A[k, j]$ 
10:      end for
11:       $A[i, k] := 0$ 
12:    end for
13:  end for
14: end function

```

观察高斯消去算法，注意到伪代码第 4, 5 行第一个内嵌循环中的 $A[k, j] := A[k, j]/A[k, k]$ 以及伪代码第 8 9 10 行双层 *for* 循环中的 $A[i, j] := A[i, j] - A[i, k] \times A[k, j]$ 都是可以进行向量化的循环。可以通过 SIMD 扩展指令对这两步进行并行优化。

2 Pthreads 算法设计

源码链接: <https://github.com/ArcanusNEO/Parallel-Programming/tree/master/4>

2.1 测试用例的确定

由于测试数据集较大, 不便于各个平台同步, 所以采用固定随机数种子为 12345687 的 mt19937 随机数生成器。经过实验发现不同规模下, 所有元素独立生成, 限制大小在 $[0, 100]$, 能够生成可以被正确消元的矩阵。

代码如下:

测试数据集生成器

```
1 uniform_real_distribution<float> dist(0, 100);
2 mt19937 mt(12345687);
3 int n;
4 istream iss(argv[1]);
5 iss >> n;
6 cout << n << endl;
7 for (int i = 1; i <= n; ++i)
8     for (int j = 1; j <= n; ++j) cout << dist(mt) << " \n"[j == n];
```

2.2 实验环境和相关配置

实验在华为鲲鹏 ARM 集群平台和本地 Arch Linux x86_64 平台完成;

华为鲲鹏 ARM 集群平台使用毕昇的 clang++ 编译器, 本地 Arch Linux x86_64 平台使用 GNU GCC 编译器;

使用 cmake 构建项目, 编译开关如下:

```
1 set(CMAKE_CXX_FLAGS_RELEASE "-O3")
2 set(THREADS_PREFER_PTHREAD_FLAG ON)
3 find_package(Threads REQUIRED)
```

2.3 算法设计

2.3.1 默认平凡算法

使用一维数组模拟矩阵, 避免改变矩阵大小时第二维不方便调整、必须设成最大值的问题, 可以减少 cache 失效;

使用 `#define matrix(i, j) arr[(i) * n + (j)]` 宏, 增强可读性;

平凡算法

```
1 #define matrix(i, j) arr[(i) * n + (j)]
2 void func(int& ans, float arr[], int n) {
3     for (int k = 0; k < n; ++k) {
4         for (int j = k + 1; j < n; ++j) matrix(k, j) = matrix(k, j) / matrix(k, k);
```

Scale	Reperat times	x86 ordinary (s)	arm ordinary (s)
8×8	100	0.000001330460	0.000000525400
16×16	50	0.000001706920	0.000001666000
32×32	50	0.000003640080	0.000007127000
64×64	20	0.000015253300	0.000037566500
128×128	15	0.000098880800	0.000231574000
256×256	10	0.000716408500	0.001820356000
512×512	10	0.006722607300	0.014974396000
1024×1024	5	0.064893815400	0.135511226000
2048×2048	3	1.400074583333	1.101775523333
4096×4096	1	10.705585484000	13.088073440000

表 1: 所有平台平凡算法结果对比

```

5     matrix(k, k) = 1.0;
6     for (int i = k + 1; i < n; ++i) {
7         for (int j = k + 1; j < n; ++j)
8             matrix(i, j) = matrix(i, j) - matrix(i, k) * matrix(k, j);
9         matrix(i, k) = 0;
10    }
11 }
12 #undef matrix
13 }

```

2.3.2 使用 Pthreads 动态创建线程并行化加速

动态创建线程 frame

```

1  #define matrix(i, j) arr[(i) * n + (j)]
2
3  #define MAX_SUB_THREAD 7
4
5  int    n;
6  float* arr;
7
8  struct thread_param_t {
9      int k, t_id;
10 };
11
12 pthread_t      thread_handle[MAX_SUB_THREAD];
13 thread_param_t thread_param[MAX_SUB_THREAD];
14
15 void* thread_func(void* param) {
16     auto p    = (thread_param_t*) param;

```

```

17     auto k    = p->k;
18     auto t_id = p->t_id;
19     int i     = k + t_id + 1;
20     for (int j = k + 1; j < n; ++j)
21         matrix(i, j) = matrix(i, j) - matrix(i, k) * matrix(k, j);
22     matrix(i, k) = 0;
23     pthread_exit(nullptr);
24 }
25
26 void func(int& ans, float arr[], int n) {
27     ::n    = n;
28     ::arr = arr;
29     for (int k = 0; k < n; ++k) {
30         for (int j = k + 1; j < n; ++j) matrix(k, j) = matrix(k, j) / matrix(k, k);
31         matrix(k, k)    = 1.0;
32         int worker_count = n - 1 - k;
33         for (int offset = 0; offset < worker_count; offset += MAX_SUB_THREAD) {
34             for (int t_id = 0, i = t_id + offset;
35                  i < worker_count && t_id < MAX_SUB_THREAD;
36                  ++t_id, i = t_id + offset) {
37                 thread_param[t_id] = {k, i};
38                 pthread_create(thread_handle + t_id, nullptr, thread_func,
39                               thread_param + t_id);
40             }
41             for (int t_id = 0, i = t_id + offset;
42                  i < worker_count && t_id < MAX_SUB_THREAD; ++t_id, i = t_id + offset)
43                 pthread_join(thread_handle[t_id], nullptr);
44         }
45     }
46 }
47 #undef matrix
48 }

```

2.3.3 使用 Pthreads 线程池和信号量同步并行化加速

线程池 + 信号量同步 + 主线程执行除法

```

1  #define matrix(i, j) arr[(i) * n + (j)]
2
3  #define MAX_SUB_THREAD 7
4
5  int n;
6  float* arr;
7
8  struct thread_param_t {
9      int t_id;
10 };
11

```

Scale	Reperat times	x86 dynamic (s)	arm dynamic (s)
8×8	100	0.000408343660	0.001069426000
16×16	50	0.001760123900	0.004892212000
32×32	50	0.007970805160	0.020191436600
64×64	20	0.031911559450	0.080842077000
128×128	15	0.130038008533	0.338396283333
256×256	10	0.515502232000	1.318216261000
512×512	10	2.105457739200	5.282467893000
1024×1024	5	8.558513793600	21.790754608000
2048×2048	3	33.120794944000	85.470935310000
4096×4096	1	139.134070615000	353.288384510000

表 2: 所有平台动态线程结果对比

```

12 sem_t      sem_main;
13 sem_t      sem_workerstart [MAX_SUB_THREAD];
14 pthread_t   handle [MAX_SUB_THREAD];
15 thread_param_t param [MAX_SUB_THREAD];
16
17 void* thread_func(void* param) {
18     auto p    = (thread_param_t*) param;
19     auto t_id = p->t_id;
20     for (int k = 0; k < n; ++k) {
21         sem_wait(sem_workerstart + t_id);
22         for (int i = k + 1 + t_id; i < n; i += MAX_SUB_THREAD) {
23             for (int j = k + 1; j < n; ++j)
24                 matrix(i, j) = matrix(i, j) - matrix(i, k) * matrix(k, j);
25             matrix(i, k) = 0;
26         }
27         sem_post(&sem_main);
28     }
29     pthread_exit(nullptr);
30 }
31
32 void func(int& ans, float arr[], int n) {
33     ::n    = n;
34     ::arr = arr;
35     sem_init(&sem_main, 0, 0);
36     for (int i = 0; i < MAX_SUB_THREAD; ++i) sem_init(sem_workerstart + i, 0, 0);
37     for (int t_id = 0; t_id < MAX_SUB_THREAD; ++t_id) {
38         param[t_id].t_id = t_id;
39         pthread_create(handle + t_id, nullptr, thread_func, param + t_id);
40     }
41     for (int k = 0; k < n; ++k) {
42         for (int j = k + 1; j < n; ++j) matrix(k, j) = matrix(k, j) / matrix(k, k);

```

Scale	Reperat times	x86 semaphore (s)	arm semaphore (s)
8×8	100	0.000220605390	0.000414891100
16×16	50	0.000319653680	0.000546852800
32×32	50	0.000513098020	0.000823440200
64×64	20	0.000966301100	0.001490467500
128×128	15	0.001879047067	0.002734098000
256×256	10	0.003608040400	0.005964593000
512×512	10	0.011124012800	0.018504859000
1024×1024	5	0.050837437000	0.077263730000
2048×2048	3	1.160066842000	0.541380500000
4096×4096	1	11.558711337000	5.892652880000

表 3: 所有平台线程池 + 信号量同步 + 主线程执行除法结果对比

```

43     matrix(k, k) = 1.0;
44     for (int t_id = 0; t_id < MAX_SUB_THREAD; ++t_id)
45         sem_post(sem_workerstart + t_id);
46     for (int t_id = 0; t_id < MAX_SUB_THREAD; ++t_id) sem_wait(&sem_main);
47 }
48 for (int t_id = 0; t_id < MAX_SUB_THREAD; ++t_id)
49     pthread_join(handle[t_id], nullptr);
50 sem_destroy(&sem_main);
51 for (int i = 0; i < MAX_SUB_THREAD; ++i) sem_destroy(sem_workerstart + i);
52 }
53 #undef matrix
54 }

```

线程池 + 信号量同步 + 工作线程执行除法

```

1  #define matrix(i, j) arr[(i) * n + (j)]
2
3  #define MAX_SUB_THREAD 7
4
5  int    n;
6  float* arr;
7
8  struct thread_param_t {
9      int t_id;
10 };
11
12 sem_t      sem_leader;
13 sem_t      sem_div[MAX_SUB_THREAD - 1];
14 sem_t      sem_elim[MAX_SUB_THREAD - 1];
15 pthread_t   handle[MAX_SUB_THREAD];
16 thread_param_t param[MAX_SUB_THREAD];

```



```

17
18 void* thread_func(void* param) {
19     auto p    = (thread_param_t*) param;
20     auto t_id = p->t_id;
21     for (int k = 0; k < n; ++k) {
22         if (t_id == 0) {
23             for (int j = k + 1; j < n; ++j)
24                 matrix(k, j) = matrix(k, j) / matrix(k, k);
25             matrix(k, k) = 1.0;
26         } else sem_wait(sem_div + t_id - 1);
27         if (t_id == 0)
28             for (int i = 0; i < MAX_SUB_THREAD - 1; ++i) sem_post(sem_div + i);
29         for (int i = k + 1 + t_id; i < n; i += MAX_SUB_THREAD) {
30             for (int j = k + 1; j < n; ++j)
31                 matrix(i, j) = matrix(i, j) - matrix(i, k) * matrix(k, j);
32             matrix(i, k) = 0.0;
33         }
34         if (t_id == 0) {
35             for (int i = 0; i < MAX_SUB_THREAD - 1; ++i) sem_wait(&sem_leader);
36             for (int i = 0; i < MAX_SUB_THREAD - 1; ++i) sem_post(sem_elim + i);
37         } else {
38             sem_post(&sem_leader);
39             sem_wait(sem_elim + t_id - 1);
40         }
41     }
42     pthread_exit(nullptr);
43 }
44
45 void func(int& ans, float arr[], int n) {
46     ::n    = n;
47     ::arr = arr;
48     sem_init(&sem_leader, 0, 0);
49     for (int i = 0; i < MAX_SUB_THREAD - 1; ++i) {
50         sem_init(sem_div + i, 0, 0);
51         sem_init(sem_elim + i, 0, 0);
52     }
53     for (int t_id = 0; t_id < MAX_SUB_THREAD; ++t_id) {
54         param[t_id].t_id = t_id;
55         pthread_create(handle + t_id, nullptr, thread_func, param + t_id);
56     }
57     for (int t_id = 0; t_id < MAX_SUB_THREAD; ++t_id)
58         pthread_join(handle[t_id], nullptr);
59     sem_destroy(&sem_leader);
60     for (int i = 0; i < MAX_SUB_THREAD - 1; ++i) {
61         sem_destroy(sem_div + i);
62         sem_destroy(sem_elim + i);
63     }
64 }
65 #undef matrix

```

Scale	Reperat times	x86 semaphore all (s)	arm semaphore all (s)
8×8	100	0.000295141480	0.000451200400
16×16	50	0.000439524060	0.000708521400
32×32	50	0.000699135760	0.001197251400
64×64	20	0.001294170900	0.001791016000
128×128	15	0.002380586333	0.004017182667
256×256	10	0.005267385000	0.009138907000
512×512	10	0.013941518400	0.020430503000
1024×1024	5	0.065091814800	0.088223814000
2048×2048	3	1.268976019000	0.604317156667
4096×4096	1	11.411054368000	5.209506790000

表 4: 所有平台线程池 + 信号量同步 + 工作线程执行除法结果对比

2.3.4 使用 Pthreads 线程池和 barrier 栅栏同步并行化加速

线程池 + 栅栏同步 + 工作线程执行除法

```

1  #define matrix(i, j) arr[(i) * n + (j)]
2
3  #define MAX_SUB_THREAD 7
4
5  int    n;
6  float* arr;
7
8  struct thread_param_t {
9      int t_id;
10 };
11
12 pthread_barrier_t barrier_div;
13 pthread_barrier_t barrier_elim;
14 pthread_t        handle[MAX_SUB_THREAD];
15 thread_param_t    param[MAX_SUB_THREAD];
16
17 void* thread_func(void* param) {
18     auto p = (thread_param_t*) param;
19     auto t_id = p->t_id;
20     for (int k = 0; k < n; ++k) {
21         if (t_id == 0) {
22             for (int j = k + 1; j < n; ++j)
23                 matrix(k, j) = matrix(k, j) / matrix(k, k);
24             matrix(k, k) = 1.0;
25         }
26
27         pthread_barrier_wait(&barrier_div);
28
29         for (int i = k + 1 + t_id; i < n; i += MAX_SUB_THREAD) {

```

Scale	Reperat times	x86 barrier (s)	arm barrier (s)
8×8	100	0.000356697420	0.000430526600
16×16	50	0.000631893380	0.000631723800
32×32	50	0.000695531360	0.000821696600
64×64	20	0.001590516150	0.001415785500
128×128	15	0.003217344067	0.003097213333
256×256	10	0.004891007500	0.005816473000
512×512	10	0.013557604400	0.017091460000
1024×1024	5	0.055628581000	0.082351834000
2048×2048	3	1.297545969000	0.503372826667
4096×4096	1	12.573134918000	5.017878950000

表 5: 所有平台线程池 + 栅栏同步 + 工作线程执行除法结果对比

```

30     for (int j = k + 1; j < n; ++j)
31         matrix(i, j) = matrix(i, j) - matrix(i, k) * matrix(k, j);
32     matrix(i, k) = 0.0;
33 }
34
35     pthread_barrier_wait(&barrier_elim);
36 }
37     pthread_exit(nullptr);
38 }
39
40 void func(int& ans, float arr[], int n) {
41     ::n = n;
42     ::arr = arr;
43     pthread_barrier_init(&barrier_div, nullptr, MAX_SUB_THREAD);
44     pthread_barrier_init(&barrier_elim, nullptr, MAX_SUB_THREAD);
45
46     for (int t_id = 0; t_id < MAX_SUB_THREAD; ++t_id) {
47         param[t_id].t_id = t_id;
48         pthread_create(handle + t_id, nullptr, thread_func, param + t_id);
49     }
50     for (int t_id = 0; t_id < MAX_SUB_THREAD; ++t_id)
51         pthread_join(handle[t_id], nullptr);
52
53     pthread_barrier_destroy(&barrier_div);
54     pthread_barrier_destroy(&barrier_elim);
55 }
56 #undef matrix
57 }

```

Scale	Reperat times	x86 ordinary (s)	x86 dynamic (s)	x86 semaphore (s)	x86 semaphore all (s)	x86 barrier (s)
8 × 8	100	0.000001330460	0.000408343660	0.000220605390	0.000295141480	0.000356697420
16 × 16	50	0.000001706920	0.001760123900	0.000319653680	0.000439524060	0.000631893380
32 × 32	50	0.000003640080	0.007970805160	0.000513098020	0.000699135760	0.000695531360
64 × 64	20	0.000015253300	0.031911559450	0.000966301100	0.001294170900	0.001590516150
128 × 128	15	0.000098880800	0.130038008533	0.001879047067	0.002380586333	0.003217344067
256 × 256	10	0.000716408500	0.515502232000	0.003608040400	0.005267385000	0.004891007500
512 × 512	10	0.006722607300	2.105457739200	0.011124012800	0.013941518400	0.013557604400
1024 × 1024	5	0.064893815400	8.558513793600	0.050837437000	0.065091814800	0.055628581000
2048 × 2048	3	1.400074583333	33.120794944000	1.160066842000	1.268976019000	1.297545969000
4096 × 4096	1	10.705585484000	139.134070615000	11.558711337000	11.411054368000	12.573134918000

表 6: x86 平台所有结果对比

Scale	Reperat times	arm ordinary (s)	arm dynamic (s)	arm semaphore (s)	arm semaphore all (s)	arm barrier (s)
8 × 8	100	0.000000525400	0.001069426000	0.000414891100	0.000451200400	0.000430526600
16 × 16	50	0.000001666000	0.004892212000	0.000546852800	0.000708521400	0.000631723800
32 × 32	50	0.000007127000	0.020191436600	0.000823440200	0.001197251400	0.000821696600
64 × 64	20	0.000037566500	0.080842077000	0.001490467500	0.001791016000	0.001415785500
128 × 128	15	0.000231574000	0.338396283333	0.002734098000	0.004017182667	0.003097213333
256 × 256	10	0.001820356000	1.318216261000	0.005964593000	0.009138907000	0.005816473000
512 × 512	10	0.014974396000	5.282467893000	0.018504859000	0.020430503000	0.017091460000
1024 × 1024	5	0.135511226000	21.790754608000	0.077263730000	0.088223814000	0.082351834000
2048 × 2048	3	1.101775523333	85.470935310000	0.541380500000	0.604317156667	0.503372826667
4096 × 4096	1	13.088073440000	353.288384510000	5.892652880000	5.209506790000	5.017878950000

表 7: arm 平台所有结果对比

3 实验结果分析

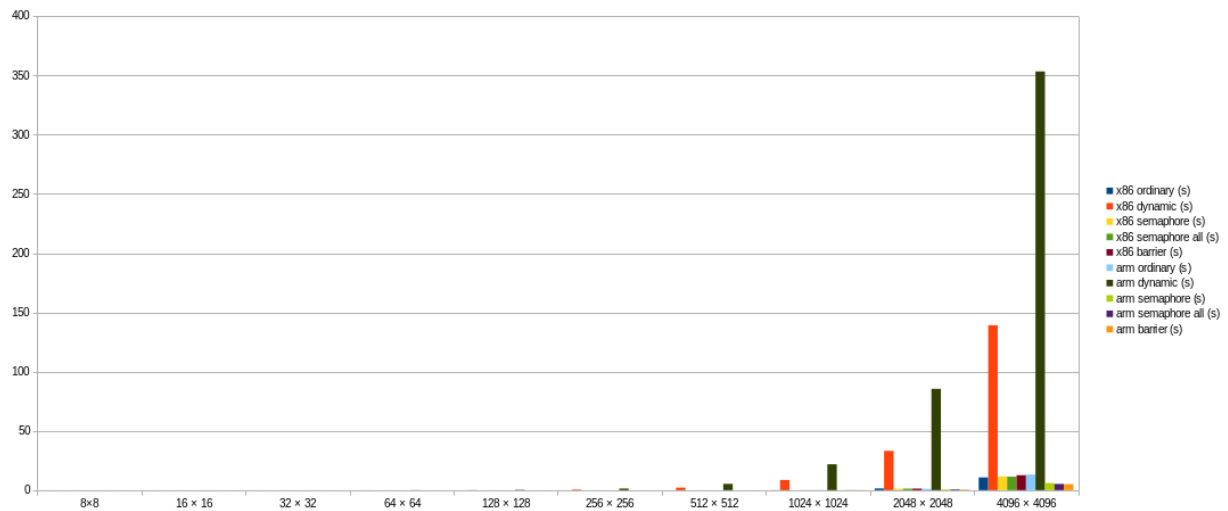


图 3.2: 所有平台所有结果对比柱状图

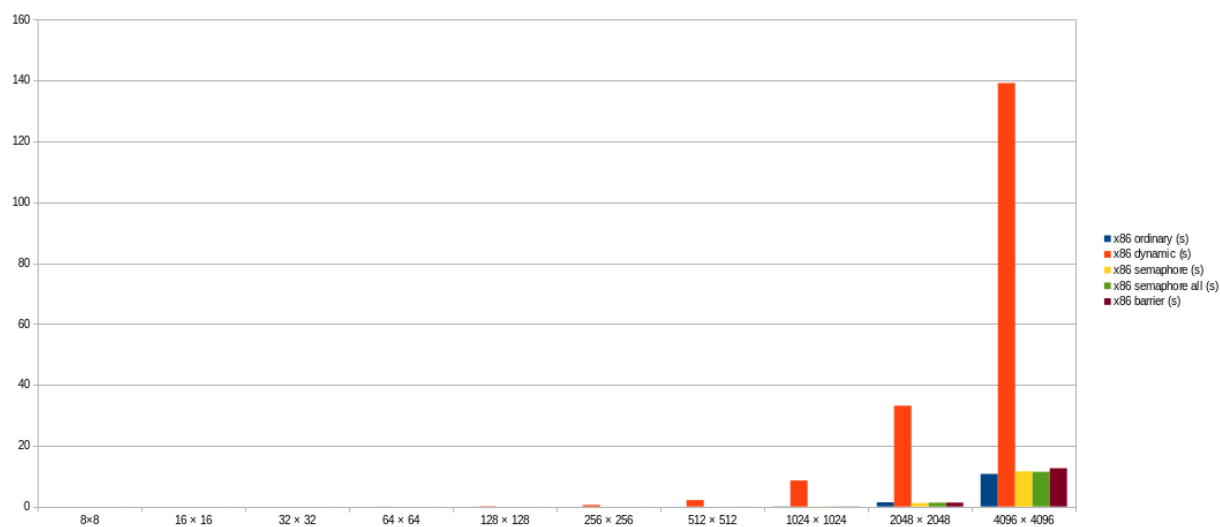


图 3.3: x86 平台所有结果对比柱状图

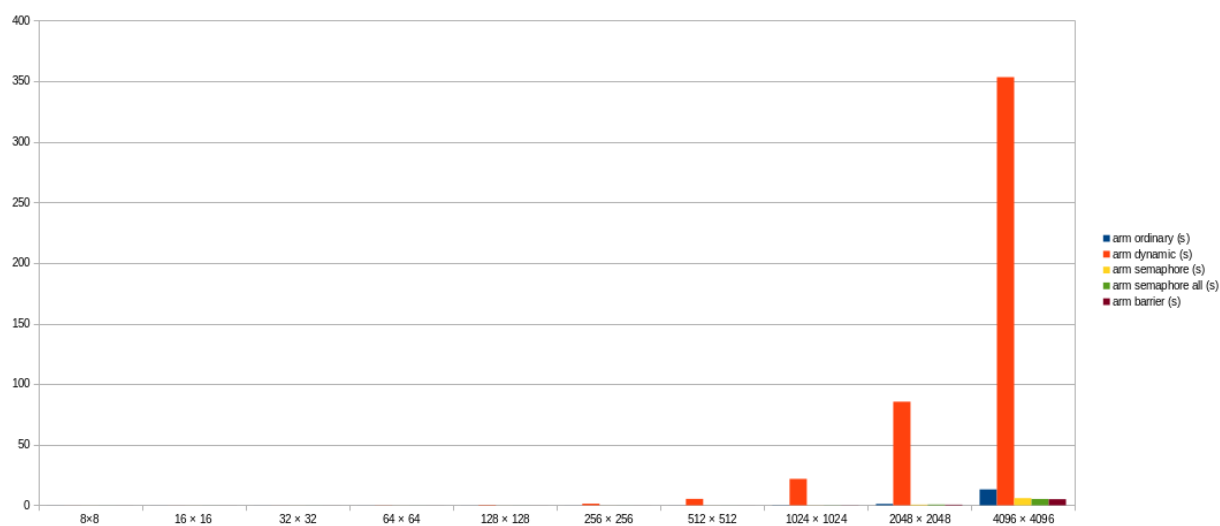


图 3.4: arm 平台所有结果对比柱状图